

Multithreading I

Jonathan R. Madsen

Department of Nuclear Engineering
Texas A&M University
College Station, TX, USA 77843
madsen_jr@tamu.edu



NUCLEAR ENGINEERING
TEXAS A & M UNIVERSITY

Outline

- 1 What is multithreading (MT)?
- 2 Geant4 MT Design
- 3 How to activate MT
- 4 How to migrate code
- 5 MT-ready Components
- 6 Summary

Overview

- Modern CPU architectures introduce **parallelism**
 - CPU frequency advances have stagnated \Rightarrow negligible advances in processing speed
 - However, the cost per microprocessor has dropped significantly \Rightarrow multiple cores instead of faster cores
- Parallelism:
 - Running tasks that are not dependent on the result of each other simultaneously on different compute nodes (cores)
 - Two core types of parallelism:
 - ① Shared-memory parallelism (multithreading)
 - ② Distributed memory parallelism (e.g. MPI)
 - Hybridization (i.e. distributed shared-memory parallelism) is possible

Threads and Processes

- Both processes and threads are independent sequences of execution
- Process
 - Has own virtual address space, executable code, open handles to system objects, security context, unique process identifier (PID), and at least one thread
- Thread
 - Entity within a process
 - All threads within a process share its virtual address space and system resources
 - Has own exception handlers, scheduling priority, thread-local storage, and unique thread identifier
- Summary: Each process has it's own memory space while threads share the memory space of their parent process

Distributed Memory Parallelism — MPI

- Multiple instances of main program (multiple PIDs)
- Each instance splits the work either by handling a specific section of the geometry (domain-decomposition) or handling a fraction of the particles (particle-decomposition) in Monte Carlo calculations
 - Particle decomposition is generally considered “embarrassingly parallel” since each process only requires a different starting random number seed
- MPI will suffer from excessive memory requirements in large problems using particle-decomposition due to the overhead of replicating the entire problem
- MPI will suffer from poor performance in many domain-decomposition cases due to communication overhead when particles leave the domain of one process and enter into the domain of another process
- Geant4 has an implementation in the extended/parallel/MPI example(s)

Shared Memory Parallelism — MT

- One instance of the main program (one unique PID)
- There is a “master” thread associated with the process, all subsequent threads are known as “worker” threads
- Subsets of the program are split into tasks, which run on threads
- Memory is shared among all the threads, although threads can allocate thread-local memory
- Geant4 has chosen MT as it's core method of parallelism
- At the core of any multithreading on UNIX systems, POSIX threads (pthreads) are used
 - Geant4 uses the pthreads library
 - Many higher level interfaces
 - OpenMP
 - Boost threads
 - Thread Building Blocks (TBB) ¹
 - C++11 threads

¹Example in [examples/extended/parallel/TBB](#)

MT vs. MPI

- Since each MPI process has it's own memory space, MPI applications generally require syncing data at some point
 - *E.g.* In a Monte Carlo application, start each MPI process with own random number seed and, at conclusion, distribute individual results to all other MPI processes — each process then has its own copy of the sum of all the results
- Since a MT process shares it's memory space, excluding thread-local memory, access to the shared memory must be governed to prevent data races (covered in MT II)
 - Thread-local memory is also generally synced, as with MPI, however, the distribution of the individual results is only sent to the “master” thread, instead of the other worker threads
 - Governing the access to the shared memory can add significant complexity to non-trivial applications

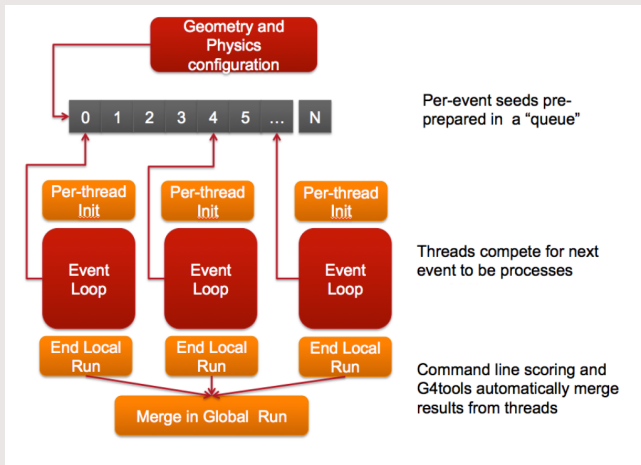
Geant4 MT

- Multithreading is introduced in Geant4 at the [G4Event²](#) level
 - In general, for performance, the higher up in the execution sequence, the better
 - Events are independent of each other — particles between events neither interact nor depend on each other
 - Each event is given its own unique random number seed
 - By reproducing the unique seeds for each event, the simulation can be reproducible in multithreaded or serial mode
- Geant4 needs back-compatibility with user code and simple approach (physicists != computer scientists)
 - We have made every effort to provide an efficient and friendly multithreaded version of Geant4 that requires a minimum amount of effort to ensure thread-safety, however, thread-safety must be kept in mind because there is only so much that can be done at the toolkit level

²This does not mean there are not thread-local G4Run instances

Geant4 MT (cont.)

General Design




Credit: Andrea Dotti, SLAC

Geant4 MT (cont.)

Simplified serial Geant4




Simplified MT Geant4



Credit: Andrea Dotti, SLAC


Geant4 MT (cont.)



Credit: Andrea Dotti, SLAC

Geant4 MT (cont.)

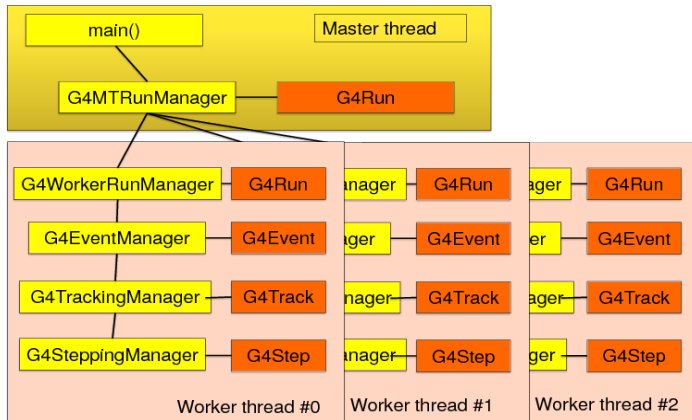
Sequential mode



Credit: Andrea Dotti, SLAC

Geant4 MT (cont.)

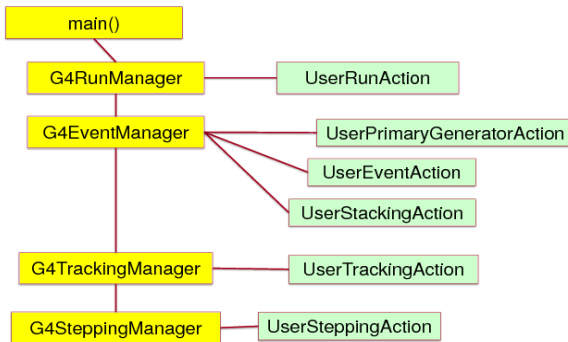
Multi-threaded mode



Credit: Andrea Dotti, SLAC

Geant4 MT (cont.)

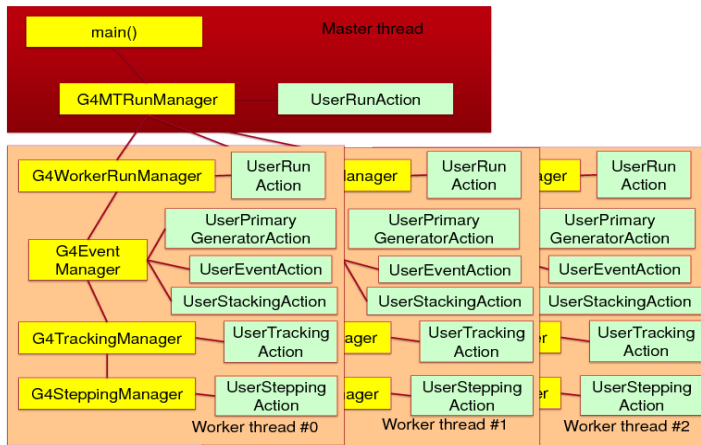
Sequential mode



Credit: Andrea Dotti, SLAC

Geant4 MT (cont.)

Multi-threaded mode



Credit: Andrea Dotti, SLAC

Shared vs. Thread-local

- To reduce memory footprint, threads must share a portion of the memory
- In general, any portion of the simulation that is static/constant³ throughout the simulation is shared
 - Geometry definitions
 - Physics data tables
 - ... etc.
- In general, for reasons explained later, any portion of the toolkit that is dynamic/changing, during the simulation is allocated to thread-local (*i.e.* “private”) memory
 - Scoring tallies
 - G4Step
 - G4Track
 - ... etc.

³ “read-only”

Compiling Geant4 with multithreading enabled (UNIX)

- POSIX threads (pthreads) is required to be installed
- `cmake -DGEANT4_BUILD_MULTITHREADED=ON [...]`
- Requires recent compiler that supports thread-local storage
 - Check cmake output:
 - `-- Performing Test HAVE_TLS`
 - `-- Performing Test HAVE_TLS -- Success`
- Geant4 applications will inherit the MT mode of the Geant4 toolkit it is built against and the preprocessor flag “**G4MULTITHREADED**” can be used in your application for specific MT-only code
 - In general, this preprocessor flag is only needed once in the `main()` of your application, which will be detailed later
 - All other Geant4 “threading-specific” tools detailed later will be ignored by the compiler when MT is turned off by the utilization of “**G4MULTITHREADED**” in the toolkit itself

Compiling Geant4 with multithreading enabled (Windows)

- Geant4 only supports multithreading on UNIX machines (Linux, Mac OS X) but can be utilized on Windows 10 with Windows Anniversary Update (Redstone) using the Windows Linux Subsystem (WLS)
 - Essentially, WLS \Leftrightarrow Ubuntu 14.04 (Trusty Tahr)
- Instructions for enabling Windows Linux Subsystem can be found online⁴, The installation procedure for Geant4 within the WLS will follow the instruction procedure for Geant4 on a Linux system
- Visualization can be enabled by installing Xming on Windows and setting the environment display variable in the WLS to the one specified when setting up Xming (e.g. DISPLAY=“:0.0”)
 - You will need to install the X11 libraries on the WLS + others
 - Ask me for a script I wrote to install all the necessary libraries to run Geant4 with MT, OpenGL, and Qt

⁴provided by Microsoft

Migration to Geant4 v10

- API has changed to smoothly allow multithreading (this is why it was a major release)
- A minimum of three modifications must be made:
 - [G4VUserDetectorConstruction](#)
 - [G4VUserActionInitialization](#)
 - [G4MTRunManager](#)
- [G4VUserDetectorConstruction](#)
 - [G4VPhysicalVolume*](#) [G4VUserDetectorConstruction::Construct\(\)](#);
 - Build geometry here except Sensitive Detectors and magnetic field
 - Called by master thread once
 - void [G4VUserDetectorConstruction::ConstructSDandField\(\)](#); [**mandatory**]
 - Build Sensitive Detectors and Magnetic fields here
 - Called by each thread

Migration to Geant4 v10 (cont.)

- Create a new class that inherits from `G4VUserActionInitialization` and implements:
 - void `G4VUserActionInitialization::Build()`
 - Instantiate user-actions for worker threads
 - Called by each thread
 - void `G4VUserActionInitialization::BuildForMaster()`
 - Instantiate user-actions for master (optional)
 - Called by master thread

Migration to Geant4 v10 (cont.)

- Add `G4MTRunManager` to `int main(int argc, char ** argv){...}`

```
int main(int argc, char** argv)
{
    ...
#ifdef G4MULTITHREADED
    // If Geant4 compiled with GEANT4_BUILD_MULTITHREADED=ON
    G4MTRunManager* runmanager = new G4MTRunManager();
    runmanager->SetNumberOfThreads(G4Threading::G4GetNumberOfCores());
#else
    // If Geant4 compiled with GEANT4_BUILD_MULTITHREADED=OFF
    G4RunManager* runmanager = new G4RunManager();
#endif

    // Detector initialization
    runmanager->SetUserInitialization(new DetectorConstruction);
    // Physics list
    runmanager->SetUserInitialization(new FTFP_BERT);
    // User-action initializations
    runmanager->SetUserInitialization(new ActionInitialization);
    ...
}
```

Scoring

- Geant4 sensitive detector, hits collections are MT ready
 - Hits objects, as well as sensitive detectors, are instantiated on worker threads
 - A keyword `G4ThreadLocal` (covered in MT II) will inform the user that a hit class is MT ready

```
// Below are sequential only (not MT compatible)
extern G4Allocator<MyHit>* MyHitAllocator; // in MyHit.hh
G4Allocator<MyHit>* MyHitAllocator = 0; // in MyHit.cc
```

```
// Below are MT-ready
extern G4ThreadLocal G4Allocator<MyHit>* MyHitAllocator; // in MyHit.hh
G4ThreadLocal G4Allocator<MyHit>* MyHitAllocator = 0; // in MyHit.cc
```

Analysis

- Geant4 analysis tools are MT-ready
- Histograms and profiles
 - Each thread owns its own copy of given histograms and profiles
 - At the end of the run, worker objects are “merged” into a single object on the master thread
 - A single file with merged histograms and profiles will be produced
- When using [G4AnalysisManager](#) with histograms, the [UserRunAction](#) class must be instantiated on **both the master and worker threads**

Analysis (cont.) — ntuples

- Each thread owns a copy of the ntuple
- ntuples are **not merged** at the end of the run — concatenate in analysis!
- Output files
 - Each thread will write own separate file, where the files names are generated automatically from the thread-id:

[fileName]_[ntupleName]_[thread-id].[extension]

e.g. run_eDep_1.txt

- When using ROOT output, the ntuple files per thread can be analyzed with use of the TChain class

Visualization

- Geant4 visualization is MT-ready
- Rendering is done by the master thread, based on event keeping settings
- Events are drawn directly from worker threads as soon as they are ready

Summary

- Geant4 supports parallelism via multithreading since Geant4 v10.0
- Multithreading is at the event-level
- MT is only available on Linux/OS X, but can be run on Windows within the Windows Linux Subsystem
- Migration to Geant4 v10+ requires changes to existing code written for Geant4 v9.6 and older